

CineBook – RAG-based Conversational AI Agent and Movie Support System

Project Concept

CineBook is an AI-powered customer support system designed for a movie ticket booking platform. It acts as a conversational AI agent that can:

- Answer questions about movies, theaters, and show timings
- Provide information about refunds, privacy, and booking policies
- Interact naturally with users through a chat interface

Unlike traditional chatbots that rely only on predefined rules, CineBook uses Retrieval-Augmented Generation (RAG) a method that combines data retrieval with LLM reasoning — to generate accurate, context-aware answer

Methodology

The methodology of the CineBook – RAG-based Conversational AI Agent focuses on integrating both structured data retrieval and unstructured semantic reasoning to enable intelligent, context-aware communication. The system's foundation lies in the Retrieval-Augmented Generation (RAG) architecture, which bridges deterministic database access with dynamic language model reasoning.

This hybrid approach ensures that every response is both accurate and contextually relevant, maintaining a natural conversational flow.

1) User Query and Intent Analysis

- The process begins when the user submits a message through the Streamlit Chat UI.
- The AI Agent receives this input and analyzes it to determine intent type — either:
- Structured Query: factual questions such as show timings or theater names.
- Unstructured Query: reasoning-based questions about refunds, policies, or account-related topics.
- This classification helps CineBook decide whether to perform a data lookup or initiate the RAG pipeline.

2) Handling Structured Queries (Direct Lookup)

- For factual queries, CineBook uses direct data lookup instead of relying on an LLM.
- The structured datasets — movies.json, theaters.json, and shows.json — contain all information related to:
 1. Movie titles, genres, and languages
 2. Theater names and locations
 3. Show timings and IDs
- During preprocessing, these files are merged into a lookup table that maps movies to their corresponding theaters and timings. The AI Agent performs instant filtering and mapping operations to fetch results.

Example:

If a user asks “Which movies are showing at StarMax Cinemas?”, CineBook searches the lookup table and returns:

“StarMax Cinemas (Velachery, Chennai): ‘Vikram’ (Tamil, Thriller) – 11:30 AM.”

This approach ensures zero hallucination and maximum accuracy for factual data retrieval.

3) Handling Unstructured Queries (RAG Pipeline)

- When a user asks reasoning-based questions — such as refunds, privacy, or user guidelines — CineBook activates the RAG (Retrieval-Augmented Generation) pipeline.
- The query text is converted into a vector embedding using SentenceTransformer (MiniLM-L6-v2), which captures its semantic meaning.
- This embedding is sent to the ChromaDB vector database, which stores precomputed embeddings from unstructured policy documents such as:
 1. refund_policy.json
 2. account_privacy_policy.json
 3. user_guidelines.json
- ChromaDB performs a similarity search using cosine distance to find the most contextually related text vectors.
- The top k (e.g., 5) relevant snippets are retrieved and used as context for response generation.

Example:

If the user asks “Can I get a refund if I miss my show?”, CineBook retrieves the rule “No refunds are applicable for missed shows” from the refund policy.

4) Response Generation Using LLM

- Once the relevant policy snippets are retrieved, the system constructs a context prompt that includes both:
 1. The user’s original query
 2. The retrieved knowledge base text
- This prompt is sent to the Groq Llama 3.1 8B model, which acts as the reasoning and generation engine.
- The LLM processes both inputs and produces a coherent, policy-grounded response, ensuring factual accuracy and fluency.

Example Response:

“You’re eligible for a refund if the cancellation is made at least two hours before showtime. Refunds are not applicable beyond that period.”

This mechanism ensures that the LLM doesn’t invent answers — it reasons strictly within the provided company policies.

5) Context Management and Session State

- CineBook maintains conversation continuity using `st.session_state` in Streamlit.
- Every user query and AI response are stored in session memory.
- This allows multi-turn conversations, enabling the system to understand follow-up questions like:

“What if I cancel only one hour before?”

This gives the user a smooth, human-like conversational experience.

6) Controlled Environment Setup

- All processes run in a Python virtual environment (`venv`) to ensure isolation and reproducibility.
- Required dependencies include:
`streamlit`, `chromadb`, `sentence-transformers`, `groq`, and `python-dotenv`.
- The `.env` file securely stores the Groq API key, keeping credentials safe and out of public code repositories.

7) Data Preparation and Embedding Creation

- Before runtime, the policy JSON files are preprocessed and chunked into smaller text segments suitable for embedding.
- Each paragraph or rule (e.g., refund conditions or data retention policies) is embedded using `SentenceTransformer` and stored in ChromaDB with unique metadata.
- Similarly, structured movie and theater data are processed into key-value mappings to ensure efficient filtering during lookup operations.
- This preprocessing enables instant data retrieval and faster RAG performance during query execution.

8) Decision Engine and Query Routing

The AI Agent's modular decision engine orchestrates all backend operations.

When a new query arrives:

- It first passes through the intent classifier (based on Python keyword detection).
- Depending on intent, it either:
- Triggers structured functions like `list_all_shows()`
- Or initiates the RAG retrieval process for unstructured queries
- The final response — either lookup-based or LLM-generated — is returned to the Streamlit UI for display.

9) Complete Workflow Summary

The entire CineBook workflow can be summarized as:

User Input → Intent Analysis → Route Selection → Retrieval (Lookup / RAG) → Response Generation → Display in Chat UI.

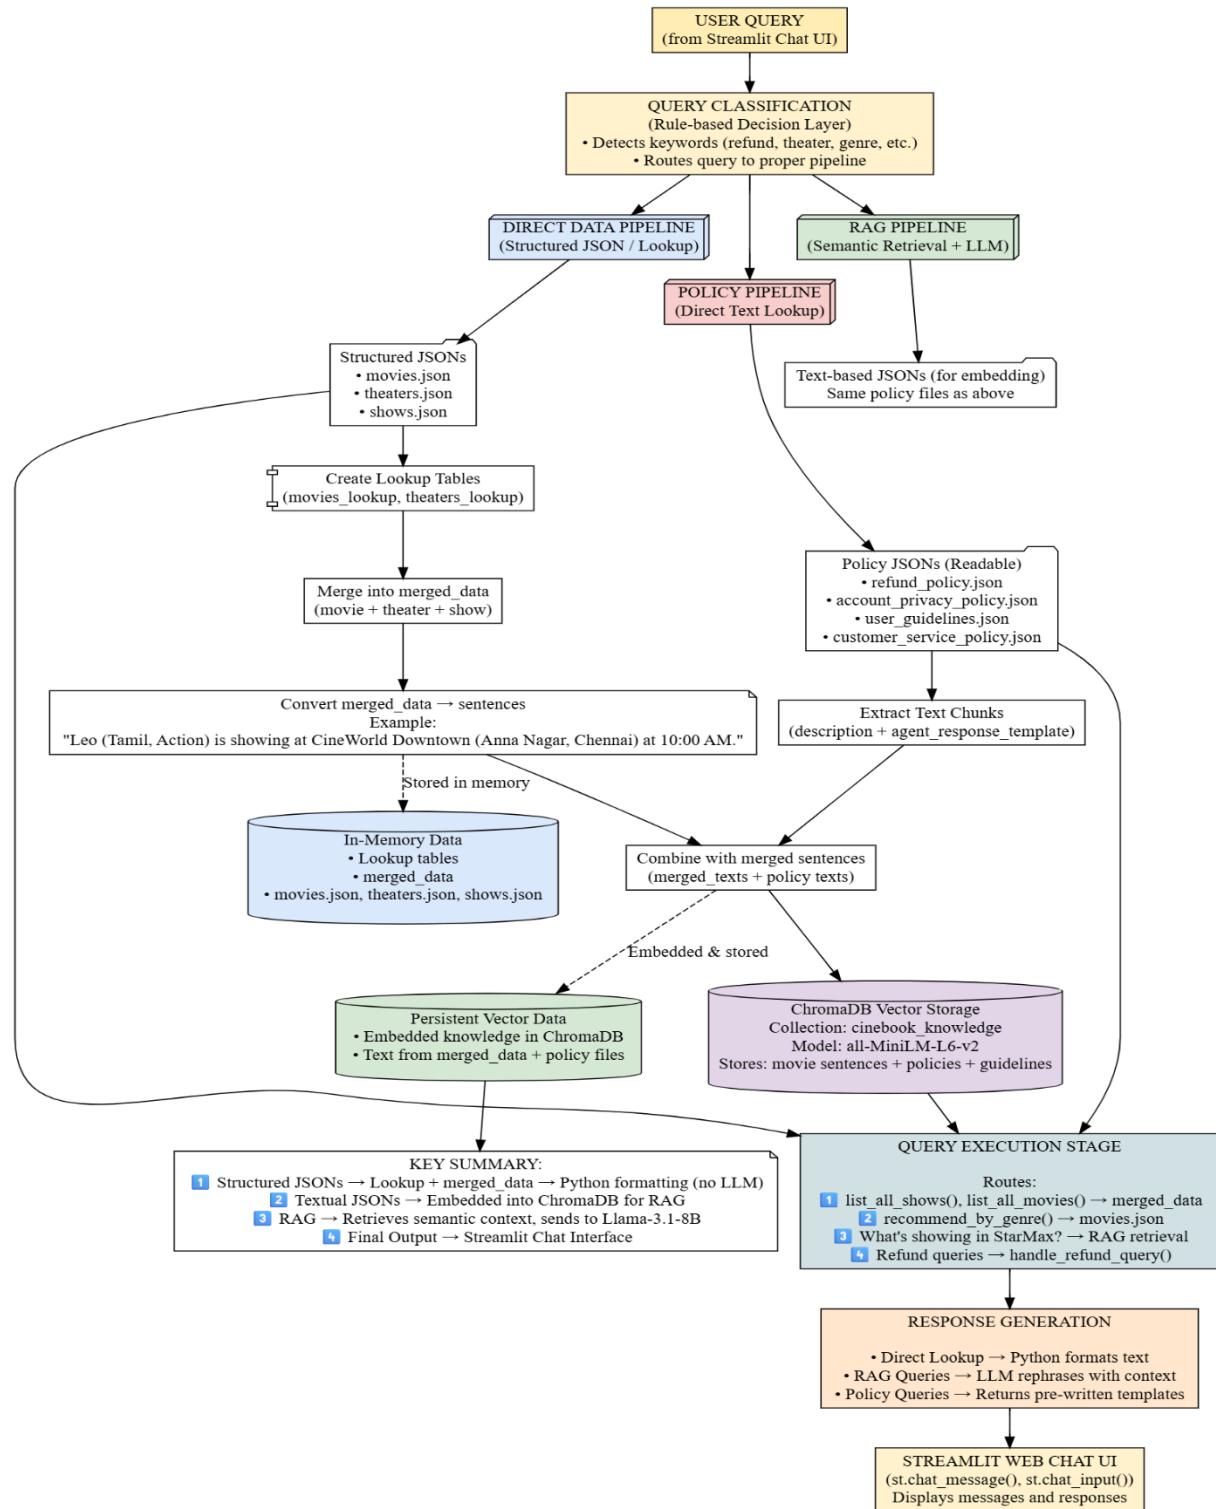
This closed-loop process allows CineBook to function as a semi-autonomous AI agent, capable of:

- Understanding the user's intent
- Retrieving relevant data or knowledge
- Reasoning contextually
- Delivering human-like, policy-compliant responses

10) Final Insight

The CineBook methodology represents a powerful fusion of symbolic reasoning (rule-based lookups) and connectionist intelligence (LLM-based generation). It demonstrates how modern AI systems can achieve domain-specific intelligence by combining structured databases with neural embeddings and generative models. This hybrid design — blending deterministic precision and semantic understanding — makes CineBook more than just a chatbot. It is a Retrieval-Augmented Conversational AI System capable of dynamic, data-driven, and policy-grounded human-like interactions.

CINIBOOK FLOWCHART



Website & interactions with user:

The screenshot shows a dark-themed web application interface. At the top, there's a navigation bar with icons for back, forward, search, and other browser functions. To the right of the bar are buttons for 'Summarize', 'Deploy', and a profile icon. The main header features a movie camera icon and the text 'CineBook — AI Movie Ticket & Policy Assistant'. Below the header, a section titled 'Welcome to CineBook AI Assistant' includes a small movie camera icon and a brief description: 'Ask about movie showtimes, get genre-based suggestions, or ask policy questions (refunds, privacy, guidelines).'. A list of 'CineBook Theaters (Chennai)' follows, with five items: CineWorld Downtown, StarMax Cinemas, Galaxy Cinemas, SilverLine Theaters, and Regal Movie Hub. A large input field at the bottom contains the placeholder 'Ask about movies, refunds, or policies...' with a red border, and a green 'Ask' button with a microphone icon to its right.

The AI assistant's response is displayed in a message bubble: 'Hello! I am an AI assistant and am here to help. Please maintain a respectful tone as per our Communication Standard. How may I assist you today?'. The user's question is shown in another message bubble: 'Yesterday, I booked a ticket for a movie 'Vikram' which was showing tomorrow. Will I get my money back, If I cancel the booking now'. The AI provides a detailed 'CineBook Refund & Cancellation Policy Summary' in a message bubble, listing cancellation rules and processing times. It also lists showtimes for the movie and a note about the cancellation window. A template message at the bottom states: 'Template: You are eligible for a refund or exchange if the cancellation is submitted at least 2 hours before the scheduled event start time.'

Hello! I am an AI assistant and am here to help. Please maintain a respectful tone as per our Communication Standard. How may I assist you today?

Yesterday, I booked a ticket for a movie 'Vikram' which was showing tomorrow. Will I get my money back, If I cancel the booking now

CineBook Refund & Cancellation Policy Summary

- ◆ Cancellation Window: 2 hours prior to the scheduled event start time (as listed on the ticket)
- ◆ Fees & Processing: Non-refundable (standard service fees are retained upon cancellation).

Processing time: 5-10 business days after approval..

- ◆ Cancellation Window: Not specified
- ◆ If provider cancels, full refund (including service fees) will be issued automatically.
- ◆ No-show: Tickets forfeited; no refund.

Found showtimes for that movie:

- StarMax Cinemas at 11:30 AM

Based on the showtime 11:30 AM on 2025-11-12, you are eligible to cancel up to 09:30 AM (2 hours before). You have approximately 791 minutes remaining to cancel.

Template: You are eligible for a refund or exchange if the cancellation is submitted at least 2 hours before the scheduled event start time.



tell me all the cinema halls with show timings



Here's the information about cinema halls with some show timings:

1. **CineWorld Downtown** in Anna Nagar:
 - **Leo** (Tamil, Action): 10:00 AM
 - **Jailer** (Tamil, Action/Drama): 02:00 PM
2. **StarMax Cinemas** in Velachery:
 - **Vikram** (Tamil, Thriller): 11:30 AM
 - **Leo** (Tamil, Action): 06:00 PM
3. **Galaxy Cinemas** in Express Avenue:
 - **Kalki 2898 AD** (Telugu, Sci-Fi): 01:00 PM
4. **SilverLine Theaters** in Royapettah:
 - **Kalki 2898 AD** (Telugu, Sci-Fi): 09:15 PM
5. **Regal Movie Hub** in Porur (Note: Show timings are for single movie - Master):
 - **Master** (Tamil, Action/Drama): 07:00 PM



suggest me a film



Let's see what's trending. I have information on the following movies:

1. **Master** (Tamil, Action/Drama)
2. **Vikram** (Tamil, Thriller)
3. **Jailer** (Tamil, Action/Drama)
4. **Kalki 2898 AD** (Telugu, Sci-Fi)
5. **Leo** (Tamil, Action)

Which genre are you interested in?



i am intrested in Sci-Fi



🎬 **Kalki 2898 AD** (Telugu, Sci-Fi) — Galaxy Cinemas at 01:00 PM 🎬 **Kalki 2898 AD** (Telugu, Sci-Fi)
— SilverLine Theaters at 09:15 PM



> 📄 Retrieved Context (for debugging/explainability)

❖ Author & Project Credits

Project Title: *CineBook – Retrieval-Augmented Conversational AI Agent & Movie Support System*

Version: 1.0

Year: 2025

👤 Author

Kavinvelavan M

LinkedIn: www.linkedin.com/in/m-kavinvelavan

GitHub: <https://github.com/mkavinvelavan>

Email: mkavinvelavan@outlook.com

❖ Tech Stack Used

- Streamlit (Frontend UI)
- ChromaDB Vector DatabaseFastAPI
- SentenceTransformer (MiniLM-L6-v2) Embeddings
- Groq API (Llama 3.1 8B Model)
- Docker
- Python
- JSON Data Sources
- Virtual Environment (venv)
- Retrieval-Augmented Generation (RAG) Architecture

🌐 Purpose of This Documentation

This documentation was created to provide a full end-to-end understanding of the CineBook system architecture, workflow, retrieval pipeline, vector indexing, and RAG methodology for professional and academic use.

🔒 Ownership

This project has been fully designed, implemented, and documented by **Kavinvelavan M**. You are welcome to explore, fork, and learn from the code.